

[Working Name] — The Standard for Trustworthy AI-Generated Analysis

Hook

Every company doing data analysis today has the same hidden liability: nobody can trace a conclusion back to the data and method that produced it. As AI agents start generating more of that analysis, that liability compounds — because now the “analyst” you can’t audit isn’t even human.

Why Now

LLMs make it possible, for the first time, to enforce standards on the *content and reasoning* of an analysis — not just its formatting or schema. CI could always check “is there a field called methodology.” It could never check “does this conclusion actually follow from this data” until now. That’s the unlock.

The Wedge: Financial/Business Reporting Teams Publishing via AI Agents

Start where provenance is already a compliance requirement and the pain is acute: finance, ops, and forecasting teams increasingly using AI agents to draft analysis and reports. These teams already need audit trails; they don’t have one for AI-generated work. This is a narrow, budget-holding buyer — not “data analysis” broadly.

How It Works

1. **Standard** — opinionated default schema for what an analysis package must contain (data, methodology, provenance, confidence). We ship the default; customers don’t design their own.
2. **Testing/CI** — every published package is validated against the Standard; non-conforming packages are flagged, not published.
3. **Team Model** — a per-customer model (RAG over their validated corpus) that lets management ask questions across all published analysis.

4. **Planner Model** — reasons across the corpus and CI flags, detects recurring gaps, and proposes changes to the Standard.
5. **HITL Review** — a human owner approves, edits, or rejects proposed Standard changes before they go live. This closes the loop and is the actual governance product, not a formality.

Moat

Not the software — a well-funded competitor or a foundation model lab can replicate a schema-and-CI product in a quarter. The moat is the proprietary, standardized, provenance-tagged corpus that only accumulates by running the pipeline over real customer usage. Software is the wedge that earns the data; the data is the asset. This is closer to how Bloomberg or Palantir defend position than how a typical SaaS tool does.

Business Model — One Primary Path

Primary: SaaS today, data-advantaged AI product tomorrow. We charge for the Standard + CI + Team Model as software from day one (real, fundable revenue). As the corpus compounds across customers, the Planner Model and pooled pattern-learning become the long-term differentiated product — sold as an upgrade tier, not a separate business.

Services (consulting-style implementation) is a funding bridge for early enterprise onboarding, not a pitched pillar — it buys us runway and design partners, not a target business model.

The Trust Answer (not just the objection)

We commit, contractually, that we train only on **structure and conformance behavior** — how packages get flagged and fixed — never on underlying business content, and never pool one customer's data into a model resold to another. Per-customer Team Models stay isolated to that customer's instance. This is decided, not open — it's the only stance that lets us sell to competitors in the same industry without an immediate veto.

Traction Plan (next 90 days)

- Ship the default Standard + CI enforcement for one document type (e.g. quarterly forecast reports) with 1 design-partner team.
- Success bar: **CI-flagged packages drop by a set % over 4–6 weeks, and 100% of**

published packages carry complete, queryable provenance — vs. 0% baseline today.

- That's the smallest provable claim: **"packages published under the Standard have complete, queryable provenance; packages published without it don't."** Everything else (Team Model, Planner Model, data moat) is sequenced after this proves out.

Market Size

The adjacent data governance/data quality software market is <cite index="1-1">projected to grow from roughly \$5.4B in 2026 to over \$24B by 2034</cite>, and the narrower data quality tools segment alone is <cite index="6-1">estimated at \$3.5B in 2026, growing to \$10.8B by 2033</cite>. We're not selling classic data governance — we're the AI-native layer for provenance on *AI-generated* analysis specifically — a wedge inside that spend, not the whole category. This is directional, not a bottoms-up SAM; real bottoms-up sizing (target account count × ACV) is the next step before an external pitch.

The Ask

Seeking one design-partner customer — a finance or ops analytics team already using AI agents for reporting — to pilot the Standard + CI + Team Model over one quarter. Capital follows proof, not the other way around: we're not opening a funding conversation until the pilot shows the flag-rate/audit-completeness result above.

Still Open (tracked, not blocking)

- Real TAM/SAM sizing for the finance/ops-analytics wedge
- Standard versioning/migration strategy as the Standard evolves
- Final governance model: who holds HITL approval authority inside a customer org